

Carnegie Mellon University

# DATABASE SYSTEMS

## More Indexes & Filters

LECTURE #09 » 15-445/645 FALL 2025 » PROF. ANDY PAVLO

# ADMINISTRIVIA

---

**Project #1** is due Sunday Sept 28<sup>th</sup> @ 11:59pm

- See Recitation Video ([@83](#))
- Office Hours: Saturday Sept 27<sup>th</sup> @ 3:00-5:00pm in GHC 5207

**Homework #3** is due Sunday Oct 5<sup>th</sup> @ 11:59pm

**Mid-Term Exam** is on Wednesday Oct 8<sup>th</sup>

- Lectures #01–11 (inclusive)
- Study guide will be released early next week.

# INDEXES VS. FILTERS

---

An index data structure of a subset of a table's attributes that are organized and/or sorted to the location of specific tuples using those attributes.  
→ Example: B+Tree

A filter is a data structure that answers set membership queries; it tells you whether a key (likely) exists in a set but not where it is located.  
→ Example: Bloom Filter

# TODAY'S AGENDA

---

Bloom Filters

Skip Lists

Tries / Radix Trees

Inverted Indexes

Vector Indexes

# BLOOM FILTERS

---

Probabilistic data structure (bitmap) that answers set membership queries.

- False negatives will never occur.
- False positives can sometimes occur.
- See [Bloom Filter Calculator](#).

**Insert(x):**

- Use  $k$  hash functions to set bits in the filter to 1.

**Lookup(x):**

- Check whether the bits are 1 for each hash function.

# BLOOM FILTERS

Insert 'RZA'

*Bloom Filter*

0	1	2	3	4	5	6	7
0	0	0	0	1	0	1	0

$$\text{hash}_1('RZA') = 2222 \% 8 = 6$$

$$\text{hash}_2('RZA') = 4444 \% 8 = 4$$

# BLOOM FILTERS

Insert 'RZA'

Insert 'GZA'

*Bloom Filter*

0	1	2	3	4	5	6	7
0	1	0	1	1	0	1	0

$$\text{hash}_1('GZA') = 5555 \% 8 = 3$$

$$\text{hash}_2('GZA') = 7777 \% 8 = 1$$

# BLOOM FILTERS

Insert 'RZA'

Insert 'GZA'

Lookup 'RZA' → *TRUE*

*Bloom Filter*

0	1	2	3	4	5	6	7
0	1	0	1	1	0	1	0

$$\text{hash}_1('RZA') = 2222 \% 8 = 6$$

$$\text{hash}_2('RZA') = 4444 \% 8 = 4$$

# BLOOM FILTERS

Insert 'RZA'

Insert 'GZA'

Lookup 'RZA' → *TRUE*

Lookup 'Raekwon' → *FALSE*

*Bloom Filter*

0	1	2	3	4	5	6	7
0	1	0	1	1	0	1	0

$$\text{hash}_1('Raekwon') = 3333 \% 8 = 5$$

$$\text{hash}_2('Raekwon') = 8899 \% 8 = 3$$

# BLOOM FILTERS

Insert 'RZA'

Insert 'GZA'

Lookup 'RZA' → *TRUE*

Lookup 'Raekwon' → *FALSE*

Lookup 'ODB' → *TRUE*

*Bloom Filter*

0	1	2	3	4	5	6	7
0	1	0	1	1	0	1	0

$$\text{hash}_1('ODB') = 6699 \% 8 = 3$$

$$\text{hash}_2('ODB') = 9966 \% 8 = 6$$

# OTHER FILTERS

---

## Counting Bloom Filter

- Supports dynamically adding and removing keys.
- Uses integers instead of bits to count the number of occurrences of a key in a set.

# OTHER FILTERS

## Counting Bloom Filter

- Supports dynamically adding and removing keys.
- Uses integers instead of bits to count the number of occurrences of a key in a set.

Carnegie  
Mellon  
University



## Cuckoo Filter

- Also supports dynamically adding and removing keys.
- Uses a Cuckoo Hash Table but stores fingerprints instead of full keys.

Carnegie  
Mellon  
University



## Succinct Range Filter (SuRF)

- Immutable compact trie that supports approximate exact matches and range filtering.

# OTHER FILTERS

Redis

## Counting Bloom Filter

- Supports dynamically adding an element.
- Uses integers instead of bits to store occurrences of a key in a set.



## Cuckoo Filter

- Also supports dynamically adding an element.
- Uses a Cuckoo Hash Table but stores keys.



## Succinct Range Filter (SURF)

- Immutable compact trie that supports matches and range filtering.

Develop with Redis

Docs → Develop with Redis → Understand Redis data types → Probabilistic → Cuckoo filter

Quick starts  
Connect  
Understand data types

- Strings
- JSON
- Lists
- Sets
- Hashes
- Sorted sets
- Streams
- Geospatial
- Blobs
- Bitmaps
- Bitfields
- Probabilistic
- HyperLogLog
- Bloom filter
- Cuckoo filter
- t-digest
- Top-K
- Count-min sketch
- Configuration
- Time series

Interact with data

Libraries and tools

Redis products

Commands

## Cuckoo filter

Cuckoo filters are a probabilistic data structure that checks for presence of an element in a set. A Cuckoo filter, just like a Bloom filter, is a probabilistic data structure in Redis Stack that enables you to check if an element is present in a set in a very fast and space efficient way, while also allowing for deletions and showing better performance than Bloom in some scenarios.

While the Bloom filter is a bit array with flipped bits at positions decided by the hash function, a Cuckoo filter is an array of buckets, storing fingerprints of the values in one of the buckets at positions decided by the two hash functions. A membership query for item x searches the possible buckets for the fingerprint of x, and returns true if an identical fingerprint is found. A cuckoo filter's fingerprint size will directly determine the false positive rate.

### Use cases

#### Targeted ad campaigns (advertising, retail)

This application answers this question: Has the user signed up for this campaign yet?

Use a Cuckoo filter for every campaign, populated with targeted users' ids. On every visit, the user id is checked against one of the Cuckoo filters.

- If yes, the user has not signed up for campaign. Show the ad.
- If the user clicks ad and signs up, remove the user id from that Cuckoo filter.
- If no, the user has signed up for that campaign. Try the next ad/Cuckoo filter.

#### Discount code/coupon validation (retail, online shops)

This application answers this question: Has this discount code/coupon been used yet?


Use a Cuckoo filter populated with all discount codes/coupons. On every try, the entered code is checked against the filter.

- If no, the coupon is not valid.
- If yes, the coupon can be valid. Check the main database. If valid, remove from Cuckoo filter as used.

# OBSERVATION

The easiest way to implement a dynamic order-preserving index is to use a sorted linked list.


All operations have to linear search.  
→ Average Cost: **O(n)**



# OBSERVATION

The easiest way to implement a dynamic order-preserving index is to use a sorted linked list.


All operations have to linear search.  
→ Average Cost: **O(n)**



# OBSERVATION

The easiest way to implement a dynamic order-preserving index is to use a sorted linked list.

All operations have to linear search.  
→ Average Cost: **O(n)**



# SKIP LISTS

Multiple levels of linked lists with extra pointers to skip over entries.

- 1<sup>st</sup> level is a sorted list of all keys.
- 2<sup>nd</sup> level links every other key
- 3<sup>rd</sup> level links every fourth key
- Each level has ½ the keys of one below it

Maintains keys in sorted order without requiring global rebalancing.

- Approximate **O(log n)** search times.

Mostly for in-memory data structures.

- Example: LSM MemTable

## Skip Lists: A Probabilistic Alternative to Balanced Trees

*Skip lists are a data structure that can be used in place of balanced trees. Skip lists use probabilistic balancing rather than strictly enforced balancing and as a result the algorithms for insertion and deletion in skip lists are less complex and significantly faster than equivalent algorithms for balanced trees.*

### William Pugh

Binary trees can be used for representing abstract data types such as dictionaries and ordered lists. They work well when the elements are inserted in a random order. Some sequences of operations, however, can lead to highly degenerate data structures that give very poor performance. If it were possible to randomly permute the list of items to be inserted, then the search time would be proportional to the input sequence. In most cases queries must be answered on-line, so randomly permuting input is impractical. *Balanced tree algorithms attempt to maintain certain balance conditions and ensure good performance.*

Skip lists are a probabilistic alternative to balanced trees. Skip lists are balanced by consulting a random number generator. Although skip lists have had worse case performance, no rigorous analysis has been done to show that their performance (much like quicksort) when the pivot element is chosen randomly. It is very unlikely a skip list data structure will be significantly slower than a balanced tree for a dictionary more than 250 elements, the cost of which will take more than 3 times the expected time is less than one in a million). Skip lists have properties similar to that of search trees built by random insertions, yet do not require a tree to be balanced.

Building a data structure probabilistically is easier than explicitly maintaining the balance. For many applications, skip lists are a more natural representation than trees, also leading to simpler algorithms. The simplicity of skip list algorithms makes them easy to understand and implement. They can constant factor speed improvements over balanced tree and self-adjusting tree algorithms. Skip lists are also very space efficient. They use only  $\log_2 n$  nodes on average, an average of  $1/2^k$  pointers per element (or even less) and do not require balance or priority information to be stored with each node.

### Skip Lists

We might need to examine every node of the list when searching a linked list (*Figure 1a*). If the list is sorted, we can stop early once we find the item. We can also point to the node two ahead of it in the list (*Figure 1b*), we have to examine no more than  $n/2 + 1$  nodes (where  $n$  is the length of the list).

Also giving every fourth node 2 pointers four ahead (*Figure 1c*) requires that no more than  $n/3 + 2$  nodes be examined. If every  $(2^{j+1})$  node has a pointer  $2^j$  nodes ahead (*Figure 1d*), then the search time is proportional to  $\log_2 n$ . This is better than  $\log_2 n$  while only doubling the number of pointers. This data structure could be used for fast searching, but insertion and deletion would be slow.

A node that has  $k$  forward pointers is called a *level k* node.

If every  $(2^{j+1})$  node has a pointer  $2^j$  nodes ahead, then levels

are roughly powers of 2: 1, 2, 4, 8, 16, 32, 64, 128, 256, 512, 1024, 2048, 4096, etc.

25% are level 2, 12.5% are level 3 and so on. What would happen if the levels of nodes were chosen randomly, but in the same way? Every node would have a pointer to the next node, instead of pointing  $2^j$  nodes ahead, points to the next node of level  $k$  or higher. Insertions or deletions would require us to change the pointers of all nodes that have been chosen randomly when the node is inserted, need never change. Some arrangements of levels would give poor execution times, but we can see that most arrangements are fine. Because these data structures are linked lists with many pointers, the skip over intermediate nodes, I named them *skip lists*.

### SKIP LIST ALGORITHMS

This section gives algorithms to search for, insert and delete elements from a skip list. The *Search* operation returns the contents of the value associated with the desired key or failure if the key is not present. The *Insert* operation inserts a new node with the specified key and value. If the key it had not already been present. The *Delete* operation deletes the specified key. It is easy to support additional operations such as *GetMin*, *GetMax*, *GetNth*, etc.

Each element is represented by a node, the level of which is chosen randomly when the node is inserted without regard to the current maximum level of the list. A *header* node has *forward* pointers, indexed 1 through  $k$ . We do not need to store the level of a node in the node. Levels are explicitly stored in the header. The *header* of a list is the maximum level currently in the list (or 1 if the list is empty). The *header* of a list has forward pointers at levels equal to or greater than the current maximum level of the list (*MaxLevel*). The forward pointers of the header at levels higher than the current maximum level of the list point to NIL.




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*

*End*




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*

*Flip a coin to decide how many levels to add the new key into.*

*End*




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*

*Flip a coin to decide how many levels to add the new key into.*

*End*




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*

*Flip a coin to decide how many levels to add the new key into.*

*End*




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*

*Flip a coin to decide how many levels to add the new key into.*

*End*




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*

*Flip a coin to decide how many levels to add the new key into.*

*End*




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*

*Flip a coin to decide how many levels to add the new key into.*

*End*




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*

*Flip a coin to decide how many levels to add the new key into.*

*End*




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*

*Flip a coin to decide how many levels to add the new key into.*

*End*




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*

*Flip a coin to decide how many levels to add the new key into.*

*End*




# SKIP LISTS: INSERT

*Insert  $K_5$*

*Levels*


*Flip a coin to decide how many levels to add the new key into.*

*End*




# SKIP LISTS: SEARCH

*Find  $K_3$*



# SKIP LISTS: SEARCH

*Find  $K_3$*




# SKIP LISTS: SEARCH

*Find  $K_3$*

*Levels*

*End*




# SKIP LISTS: SEARCH

*Find  $K_3$*

*Levels*

*End*



# SKIP LISTS: DELETE

---

First logically remove a key from the index by setting a flag to tell threads to ignore.


Then physically remove the key once we know that no other thread is holding the reference.

# SKIP LISTS: DELETE

*Delete  $K_5$*

*Levels*

*End*




# SKIP LISTS: DELETE

*Delete  $K_5$*

*Levels*

*End*




# SKIP LISTS: DELETE

*Delete  $K_5$*

*Levels*

*End*




# SKIP LISTS: DELETE

*Delete  $K_5$*

*Levels*

*End*




# SKIP LISTS: DELETE

*Delete  $K_5$*

*Levels*

*End*




# SKIP LISTS: DELETE

*Delete  $K_5$*

*Levels*

*End*




# SKIP LISTS: DELETE

*Delete  $K_5$*

*Levels*

*End*




# SKIP LISTS: DELETE

*Delete  $K_5$*

*Levels*

*End*



# SKIP LISTS

---

## Advantages:

- May use less memory than a B+Tree, if you do not include reverse pointers.
- Insertions and deletions do not require rebalancing.

## Disadvantages:

- Not disk/cache friendly because they do not optimize locality of references.
- Reverse search is non-trivial.

# OBSERVATION

---

The inner node keys in a B+Tree cannot tell you whether a key exists in the index. You must always traverse to the leaf node.

This means that you could have (at least) one buffer pool page miss per level in the tree just to find out a key does not exist.

# TRIE INDEX

Use a digital representation of keys to examine prefixes one-by-one.

→ aka *Digital Search Tree, Prefix Tree*.

Shape depends on keys and lengths.

→ Does not depend on existing keys or insertion order.


→ Does not require rebalancing operations.

All operations have **O(k)** complexity where **k** is the length of the key.

→ Path to a leaf node represents a key.

→ Keys are stored implicitly and can be reconstructed from paths.

**Keys:** HELLO, HAT, HAVE



# TRIE INDEX

Use a digital representation of keys to examine prefixes one-by-one.

→ aka *Digital Search Tree, Prefix Tree*.

Shape depends on keys and lengths.

→ Does not depend on existing keys or insertion order.


→ Does not require rebalancing operations.

All operations have **O( $k$ )** complexity where  **$k$**  is the length of the key.

→ Path to a leaf node represents a key.

→ Keys are stored implicitly and can be reconstructed from paths.

Keys: HELLO, HAT, HAVE



# TRIE KEY SPAN

---

The span of a trie level is the number of bits that each partial key / digit represents.


- If the digit exists in the corpus, then store a pointer to the next level in the trie branch. Otherwise, store null.

This determines the fan-out of each node and the physical height of the tree.

- $n$ -way Trie = Fan-Out of  $n$

# TRIE KEY SPAN

## 1-bit Span Trie



Keys: K10, K25, K31


K10 → 000000000 00001010

K25 → 000000000 00011001

K31 → 000000000 00011111

# TRIE KEY SPAN

## 1-bit Span Trie




Keys: K10, K25, K31

K10 →	0 0 0 0 0 0 0 0	0 0 0 0 1 0 1 0
K25 →	0 0 0 0 0 0 0 0	0 0 0 1 1 0 0 1
K31 →	0 0 0 0 0 0 0 0	0 0 0 1 1 1 1 1

# TRIE KEY SPAN

## *1-bit Span Trie*




**Keys:** K10, K25, K31

K10→ 00000000 00001010  
K25→ 00000000 00011001  
K31→ 00000000 00011111

# TRIE KEY SPAN

## 1-bit Span Trie




Keys: K10, K25, K31

K10 → 000000000 00001010  
 K25 → 000000000 00011001  
 K31 → 000000000 00011111

# TRIE KEY SPAN

## 1-bit Span Trie




Keys: K10, K25, K31

K10 → 0 0 0 0 0 0 0 0	0 0 0 0 1 0 1 0
K25 → 0 0 0 0 0 0 0 0	0 0 0 1 1 0 0 1
K31 → 0 0 0 0 0 0 0 0	0 0 0 1 1 1 1 1



# TRIE KEY SPAN

## 1-bit Span Trie




Keys: K10, K25, K31

K10 →	0 0 0 0 0 0 0 0	0 0 0 0 1 0 1 0
K25 →	0 0 0 0 0 0 0 0	0 0 0 1 1 0 0 1
K31 →	0 0 0 0 0 0 0 0	0 0 0 1 1 1 1 1

# TRIE KEY SPAN

## 1-bit Span Trie




Keys: K10, K25, K31

K10 →	0 0 0 0 0 0 0 0	0 0 0 0 1 0 1 0
K25 →	0 0 0 0 0 0 0 0	0 0 0 1 1 0 0 1
K31 →	0 0 0 0 0 0 0 0	0 0 0 1 1 1 1 1

# TRIE KEY SPAN

## 1-bit Span Trie



Keys: K10, K25, K31


K10 → 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0

K25 → 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 1

K31 → 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1

# TRIE KEY SPAN

## 1-bit Span Trie




Keys: K10, K25, K31

K10 →	0 0 0 0 0 0 0 0	0 0 0 0 1 0 1 0
K25 →	0 0 0 0 0 0 0 0	0 0 0 1 1 0 0 1
K31 →	0 0 0 0 0 0 0 0	0 0 0 1 1 1 1 1

# RADIX TREE

## 1-bit Span Radix Tree




Vertically compressed trie that compacts nodes with a single child.  
→ aka Patricia Trees (2-way tries)

Can produce false positives, so the DBMS always checks the original tuple to see whether a key matches.




# RADIX TREE: MODIFICATIONS

*Insert HAIR*




# RADIX TREE: MODIFICATIONS




*Insert HAIR*

# RADIX TREE: MODIFICATIONS




*Insert HAIR*  
*Delete HAT*

# RADIX TREE: MODIFICATIONS



*Insert HAIR*  
*Delete HAT*

# RADIX TREE: MODIFICATIONS




*Insert HAIR*

*Delete HAT*

*Delete HAVE*

# RADIX TREE: MODIFICATIONS




*Insert HAIR*

*Delete HAT*

*Delete HAVE*

# RADIX TREE: MODIFICATIONS




*Insert HAIR*

*Delete HAT*

*Delete HAVE*

# RADIX TREE: MODIFICATIONS



*Insert HAIR*

*Delete HAT*

*Delete HAVE*

# TRIE VARIANTS

---

## Judy Arrays (HP)

→ 256-way radix tree. First known radix tree that supports adaptive node representation.

## ART Index (HyPer)

→ 256-way radix tree that also adapts node types based on its population.

## Masstree (Silo)

→ B+ Tree that uses tries inside each node for its key/value mapping.

## SuRF (CMU)

→ Compressed static trie that supports point and range queries.

# OBSERVATION

The indexes that we've discussed are useful for "point" and "range" queries:

- Find all customers in the 15217 zipcode.
- Find all orders between June 2024 and September 2024.

They are not good at keyword searches:

- Example: Find all Wikipedia articles that contain the word "Pavlo"

**revisions(id, content, ...)**

<b>id</b>	<b>content</b>
11	Wu-Tang Clan is an American hip hop musical collective formed in Staten Island, New York City, in 1992...
22	Carnegie Mellon University (CMU) is a private research university in Pittsburgh, Pennsylvania. The institution was established in 1900 by Andrew Carnegie...
33	In computing, a database is an organized collection of data or a type of data store based on the use of a database management system (DBMS), the software...
44	Andrew Pavlo, best known as Andy Pavlo, is an associate professor of Computer Science at Carnegie Mellon University. He conducts research on database...

```
CREATE INDEX idx_rev_cntnt
ON revisions (content);
```

```
SELECT pageID FROM revisions
WHERE content LIKE '%Pavlo%';
```

# INVERTED INDEX

An **inverted index** stores a mapping of terms to records that contain those terms in the target attribute.

- Sometimes called a *full-text search index*.
- Originally called a **concordance** (1200s).

Many major DBMSs support these natively. But there are also specialized DBMSs and libraries.

*Term / Frequency*

**revisions(id, content,...)**

id	content
11	Wu-Tang Clan is an American hip hop musical collective formed in Staten Island, New York City, in 1992...
22	Carnegie Mellon University (CMU) is a private research university in Pittsburgh, Pennsylvania. The institution was established in 1900 by Andrew Carnegie...
33	In computing, a database is an organized collection of data or a type of data store based on the use of a database management system (DBMS), the software...
44	Andrew Pavlo, best known as Andy Pavlo, is an associate professor of Computer Science at Carnegie Mellon University. He conducts research on database...

*Dictionary*

Wu-Tang | 2

Carnegie | 3

Database | 2

*Posting Lists*

11	44		
----	----	--	--

22	33	44	
----	----	----	--

33	44		
----	----	--	--

⋮

# INVERTED INDEX

An **inverted index** stores a mapping of terms to records that contain those terms in the target attribute.

- Sometimes called a *full-text search index*.
- Originally called a **concordance** (1200s).

Many major DBMSs support these natively. But there are also specialized DBMSs and libraries.



elasticsearch

Xapian



OpenSearch

QUICKWIT

vespa



Sphinx

ParadeDB

splunk>

*Term / Frequency*

**revisions(id, content, ...)**

id	content
11	Wu-Tang Clan is an American hip hop musical collective formed in Staten Island, New York City, in 1992...
22	Carnegie Mellon University (CMU) is a private research university in Pittsburgh, Pennsylvania. The institution was established in 1900 by Andrew Carnegie...
33	In computing, a database is an organized collection of data or a type of data store based on the use of a database management system (DBMS), the software...
44	Andrew Pavlo, best known as Andy Pavlo, is an associate professor of Computer Science at Carnegie Mellon University. He conducts research on database...

*Dictionary*

Wu-Tang | 2

Carnegie | 3

Database | 2

*Posting Lists*

11	44		
----	----	--	--

22	33	44	
----	----	----	--

33	44		
----	----	--	--

# INVERTED INDEX

An **inverted index** stores a mapping of terms to records that contain those terms in the target attribute.

- Sometimes called a *full-text search index*.
- Originally called a **concordance** (1200s).

Many major DBMSs support these natively. But there are also specialized DBMSs and libraries.



elasticsearch

Xapian



OpenSearch

QUICKWIT

vespa



Sphinx

ParadeDB

splunk>

*Term / Frequency*

**revisions(id, content, ...)**

id	content
11	Wu-Tang Clan is an American hip hop musical collective formed in Staten Island, New York City, in 1992...
22	Carnegie Mellon University (CMU) is a private research university in Pittsburgh, Pennsylvania. The institution was established in 1900 by Andrew Carnegie...
33	In computing, a database is an organized collection of data or a type of data store based on the use of a database management system (DBMS), the software...
44	Andrew Pavlo, best known as Andy Pavlo, is an associate professor of Computer Science at Carnegie Mellon University. He conducts research on database...

*Dictionary*

Wu-Tang | 2

Carnegie | 3

Database | 2

*Posting Lists*

11	44		
----	----	--	--

22	33	44	
----	----	----	--

33	44		
----	----	--	--

# INVERTED INDEX: LUCENE


Uses a **Finite State Transducer** for determining offset of terms in dictionary.

Incrementally create dictionary segments and then merge them in the background.

- Uses **compression methods** we previously discussed (e.g., delta, bit packing).
- Also supports precomputed aggregations for terms and occurrences.

Dictionary

1	BR
2	BRAV
3	PAV
4	PLA



# INVERTED INDEX: LUCENE


Uses a Finite State Transducer for determining offset of terms in dictionary.

Incrementally create dictionary segments and then merge them in the background.

- Uses compression methods we previously discussed (e.g., delta, bit packing).
- Also supports precomputed aggregations for terms and occurrences.

Dictionary	
1	BR
2	BRAV
3	PAV
4	PLA

*Offset = 0*



# INVERTED INDEX: LUCENE

Uses a Finite State Transducer for determining offset of terms in dictionary.

Incrementally create dictionary segments and then merge them in the background.


- Uses compression methods we previously discussed (e.g., delta, bit packing).
- Also supports precomputed aggregations for terms and occurrences.

*Dictionary*

1	BR
2	BRAV
3	PAV
4	PLA

*Offset= 2*

*Find PAV*



# INVERTED INDEX: LUCENE

Uses a **Finite State Transducer** for determining offset of terms in dictionary.

Incrementally create dictionary segments and then merge them in the background.


- Uses **compression methods** we previously discussed (e.g., delta, bit packing).
- Also supports precomputed aggregations for terms and occurrences.

Dictionary

1	BR
2	BRAV
3	PAV
4	PLA

*Offset = 3*

Find PAV




# INVERTED INDEX: LUCENE


Uses a Finite State Transducer for determining offset of terms in dictionary.

Incrementally create dictionary segments and then merge them in the background.

- Uses compression methods we previously discussed (e.g., delta, bit packing).
- Also supports precomputed aggregations for terms and occurrences.



*Find PAV*




# INVERTED INDEX: POSTGRESQL

PostgreSQL's **Generalized Inverted Index** (GIN) uses a B+Tree for the term dictionary that map to a posting list data structure.

Posting list contents varies depending on number of records per term:

- **Few**: Sorted list of record ids.
- **Many**: Another B+Tree of record ids.

Uses a separate "pending list" log to avoid incremental updates.




# INVERTED INDEX: POSTGRESQL

PostgreSQL's Generalized Inverted Index (GIN) uses a B+Tree for the term dictionary that map to a posting list data structure.

Posting list contents varies depending on number of records per term:

- **Few**: Sorted list of record ids.
- **Many**: Another B+Tree of record ids.

Uses a separate "pending list" log to avoid incremental updates.



# INVERTED INDEX: ENHANCEMENTS

**Rankings:** The DBMS can rank search results based on the frequency of terms in each record relative to other records.

→ Examples: [TF-IDF](#), [BM25](#)

# INVERTED INDEX: ENHANCEMENTS

**Rankings:** The DBMS can rank search results based on the frequency of terms in each record relative to other records.


→ Examples: [TF-IDF](#), [BM25](#)

**Tokenizers:** Split terms into *n*-grams to support fuzzy text searches and autocomplete ("did you mean").

→ Examples: Elastic N-gram, pg\_trgm

```
SELECT pageID FROM revisions
WHERE content LIKE 'the';
```

```
SELECT pageID FROM revisions
WHERE content LIKE 'Wu-Tang';
```



# OBSERVATION

---

Inverted indexes search data based on its contents with some term tweaking based on linguistic models.

→ Example: Normalization ("Wu-Tang" matches "Wu Tang").

Instead of searching for records containing keywords (e.g., "**Wu-Tang**"), an application may want search for records that are semantically similar to topics (e.g., "hip-hop groups with songs about slinging").

# SEMANTIC SIMILARITY SEARCH

**Album**(id, name, year, lyrics)

id	name	year	lyrics
Id1	<u>Enter the Wu-Tang</u>	1993	<text>
Id2	<u>Run the Jewels 2</u>	2015	<text>
Id3	<u>Liquid Swords</u>	1995	<text>
Id4	<u>We Got It from Here</u>	2016	<text>



OpenAI



Hugging Face

*Transformer*

*Embeddings*

Id1 → [0.32, 0.78, 0.30, ...]

Id2 → [0.99, 0.19, 0.81, ...]

Id3 → [0.01, 0.18, 0.85, ...]

Id4 → [0.19, 0.82, 0.24, ...]

⋮



# SEMANTIC SIMILARITY SEARCH

**Album(id, name, year, lyrics)**

id	name	year	lyrics
Id1	Enter the Wu-Tang	1993	<text>
Id2	Run the Jewels 2	2015	<text>
Id3	Liquid Swords	1995	<text>
Id4	We Got It from Here	2016	<text>



OpenAI



Hugging Face

*Transformer*

*Embeddings*

Id1 → [0.32, 0.78, 0.30, ...]
Id2 → [0.99, 0.19, 0.81, ...]
Id3 → [0.01, 0.18, 0.85, ...]
Id4 → [0.19, 0.82, 0.24, ...]
⋮

*Query*

Find albums with lyrics about  
**running from the police**

[0.02, 0.10, 0.24, ...]

*Vector Index*

# SEMANTIC SIMILARITY SEARCH

**Album(id, name, year, lyrics)**

id	name	year	lyrics
Id1	Enter the Wu-Tang	1993	<text>
Id2	Run the Jewels 2	2015	<text>
Id3	Liquid Swords	1995	<text>
Id4	We Got It from Here	2016	<text>



**Transformer**

**Query**

Find albums with lyrics about  
**running from the police**

**Embeddings**

Id1 → [0.32, 0.78, 0.30, ...]
Id2 → [0.99, 0.19, 0.81, ...]
Id3 → [0.01, 0.18, 0.85, ...]
Id4 → [0.19, 0.82, 0.24, ...]
⋮

[0.02, 0.10, 0.24, ...]

**Ranked List of Ids**

**Vector Index**

# SEMANTIC SIMILARITY SEARCH

**Album**(id, name, year, lyrics)

id	name	year	lyrics
Id1	Enter the Wu-Tang	1993	<text>
Id2	Run the Jewels 2	2015	<text>
Id3	Liquid Swords	1995	<text>
Id4	We Got It from Here	2016	<text>



**OpenAI**




**Hugging Face**

*Transformer*

*Embeddings*

Id1 → [0.32, 0.78, 0.30, ...]
Id2 → [0.99, 0.19, 0.81, ...]
Id3 → [0.01, 0.18, 0.85, ...]
Id4 → [0.19, 0.82, 0.24, ...]
⋮




# SEMANTIC SIMILARITY SEARCH

**Album**(id, name, year, lyrics)

id	name	year	lyrics
Id1	Enter the Wu-Tang	1993	<text>
Id2	Run the Jewels 2	2015	<text>
Id3	Liquid Swords	1995	<text>
Id4	We Got It from Here	2016	<text>

## Query

Find albums with lyrics about  
**running from the police**  
 and released after **2005**



# VECTOR INDEXES

Specialized data structures to perform nearest-neighbor searches on embeddings (i.e., one-dimensional arrays of floating-point numbers).

- May also need to filter data before / after vector searches.
- Result correctness depends on whether the result "feels right".

Two approaches:

- Inverted Indexes
- Graphs



turbopuffer <(°0°)>



# VECTOR INDEXES

Specialized data structures to perform nearest-neighbor searches on embeddings (i.e., one-dimensional arrays of floating-point numbers).

- May also need to filter data before / after vector searches.
- Result correctness depends on whether the result "feels right".

Two approaches:

- Inverted Indexes
- Graphs



turbopuffer <(°0°)>



# VECTOR SEARCH: INVERTED INDEXES

Partition vectors into smaller groups using a clustering algorithm and then build an inverted index that maps cluster centroids to records.

- Preprocess / quantize vectors to reduce dimensionality.
- Some implementations support incremental updates.
- Examples: [IVFFlat](#), [pgVector](#)

To find a match, use same clustering algorithm to map into a group, then scan that group's vectors.

- Also check nearby groups to improve accuracy.

# VECTOR SEARCH: INVERTED INDEXES

Compute k-means clustering on embeddings to find centroids.

Assign each embedding to the cluster with its closest centroid.

## *Embeddings*

Id1 → [0.32, 0.78, 0.30, ...]
Id2 → [0.99, 0.19, 0.81, ...]
Id3 → [0.01, 0.18, 0.85, ...]
Id4 → [0.19, 0.82, 0.24, ...]
⋮

# VECTOR SEARCH: INVERTED INDEXES

Compute k-means clustering on embeddings to find centroids.


Assign each embedding to the cluster with its closest centroid.

## Embeddings

Id1	→	[0.32, 0.78, 0.30, ...]
Id2	→	[0.99, 0.19, 0.81, ...]
Id3	→	[0.01, 0.18, 0.85, ...]
Id4	→	[0.19, 0.82, 0.24, ...]
		⋮



## Clusters



# VECTOR SEARCH: INVERTED INDEXES

Compute k-means clustering on embeddings to find centroids.


Assign each embedding to the cluster with its closest centroid.

## *Embeddings*

```
Id1 → [0.32, 0.78, 0.30, ...]  
Id2 → [0.99, 0.19, 0.81, ...]  
Id3 → [0.01, 0.18, 0.85, ...]  
Id4 → [0.19, 0.82, 0.24, ...]  
⋮
```



## *Clusters*



# VECTOR SEARCH: INVERTED INDEXES


Compute k-means clustering on embeddings to find centroids.

Assign each embedding to the cluster with its closest centroid.

## Embeddings

Id1	→	[0.32, 0.78, 0.30, ...]
Id2	→	[0.99, 0.19, 0.81, ...]
Id3	→	[0.01, 0.18, 0.85, ...]
Id4	→	[0.19, 0.82, 0.24, ...]
		⋮

## Clusters



# VECTOR SEARCH: INVERTED INDEXES

Compute k-means clustering on embeddings to find centroids.

Assign each embedding to the cluster with its closest centroid.


Build an inverted index mapping centroids to a posting list of the records in each cluster.

To find matches, compute centroid of closest cluster on search embedding and then evaluate posting list

*Dictionary      Posting Lists*

centroid1	→	Id1 Id4 Id19 Id33
centroid2	→	Id2 Id10 Id24 Id20 Id19
centroid3	→	Id3 Id6
centroid4	→	Id4 Id8 Id16 Id32 Id64
⋮		

*Clusters*




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its **n** nearest neighbors.  
→ Can use multiple levels of graphs ([HNSW](#))  
→ Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

*Vector Graph*




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

*Vector Graph*




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

*Vector Graph*




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

*Vector Graph*




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

*Vector Graph*




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

*Vector Graph*




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its **n** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

*Hierarchical Graph*




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

## Hierarchical Graph




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

## Hierarchical Graph




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

*Hierarchical Graph*




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

*Hierarchical Graph*




# VECTOR SEARCH: GRAPH

Build a graph where each node represents a vector and it has edges to (at most) its ***n*** nearest neighbors.

- Can use multiple levels of graphs ([HNSW](#))
- Examples: [Faiss](#), [hnswlib](#), [DiskANN](#)

To find a match for a given vector, enter the graph and then greedily choose the next edge that moves closer to that vector.

## Hierarchical Graph



# PARTIAL INDEXES

Create an index on a subset of the entire table. This potentially reduces its size and the amount of overhead to maintain it.

```
CREATE INDEX idx_foo  
    ON foo (a, b)  
WHERE c = 'WuTang';
```

One common use case is to partition indexes by date ranges.

→ Create a separate index per month, year.

# PARTIAL INDEXES

Create an index on a subset of the entire table. This potentially reduces its size and the amount of overhead to maintain it.

One common use case is to partition indexes by date ranges.

→ Create a separate index per month, year.

```
CREATE INDEX idx_foo  
    ON foo (a, b)  
    WHERE c = 'WuTang';
```

```
SELECT b FROM foo  
    WHERE a = 123  
        AND c = 'WuTang';
```

# PARTIAL INDEXES

Create an index on a subset of the entire table. This potentially reduces its size and the amount of overhead to maintain it.

One common use case is to partition indexes by date ranges.

→ Create a separate index per month, year.

```
CREATE INDEX idx_foo  
    ON foo(a, b)  
    WHERE c = 'WuTang';
```

```
SELECT b FROM foo  
WHERE a = 123
```

# INDEX INCLUDE COLUMNS

Embed additional columns in indexes to support index-only queries.

These extra columns are only stored in the leaf nodes and are not part of the search key.

```
CREATE INDEX idx_foo  
    ON foo (a, b)  
    INCLUDE (c);
```

```
SELECT b FROM foo  
WHERE a = 123  
AND c = 'WuTang';
```

# INDEX INCLUDE COLUMNS

Embed additional columns in indexes to support index-only queries.

These extra columns are only stored in the leaf nodes and are not part of the search key.

```
CREATE INDEX idx_foo  
    ON foo (a, b)  
    INCLUDE (c);
```

```
SELECT b FROM foo  
WHERE a = 123  
    AND c = 'WuTang';
```

# INDEX INCLUDE COLUMNS

Embed additional columns in indexes  
to support index-only queries.

These extra columns are only stored  
in the leaf nodes and are not part of  
the search key.

```
CREATE INDEX idx_foo  
    ON foo(a, b)  
    INCLUDE (c);
```

```
SELECT b FROM foo  
WHERE a = 123  
    AND c = 'WuTang';
```

# INDEX INCLUDE COLUMNS

Embed additional columns in indexes  
to support index-only queries.

These extra columns are only stored  
in the leaf nodes and are not part of  
the search key.

```
CREATE INDEX idx_foo  
    ON foo(a, b)  
    INCLUDE (c)
```

```
SELECT b FROM foo  
WHERE a = 123  
    AND c = 'WuTang' ;
```

# INDEX INCLUDE COLUMNS

Embed additional columns in indexes to support index-only queries.

These extra columns are only stored in the leaf nodes and are not part of the search key.

If all attributes a query needs are available in an index, then the DBMS does not need to retrieve the tuple.  
→ AKA covering index or index-only scans.

```
CREATE INDEX idx_foo  
    ON foo(a, b)  
    INCLUDE (c)
```

```
SELECT b FROM foo  
WHERE a = 123  
    AND c = 'WuTang' ;
```

# CONCLUSION

---

We will see filters again this semester.

B+ Trees are still the way to go for tree indexes.

Inverted indexes are covered in [CMU 11-442](#).

We did not discuss geo-spatial tree indexes:

- Examples: R-Tree, Quad-Tree, KD-Tree
- This is covered in [CMU 15-826](#).

# NEXT CLASS

---

How to make indexes thread-safe!